

# Dynamic Sporadic Server

Moiz Zaheer Malik

B.Eng. Electronic Engineering

Hochschule Hamm-Lippstadt

Lippstadt, Germany

moiz-zaheer.malik@stud.hshl.de

**Abstract**—In real-time systems, where tasks have different levels of critical importance, it is essential to serve aperiodic (irregular, event-driven) tasks while ensuring that the deadlines of high-priority periodic tasks are not violated. The Dynamic Sporadic Server (DSS) is a scheduling method designed for Earliest Deadline First (EDF) systems that addresses this problem.

DSS is defined by a period  $T_s$  and a budget  $C_s$ , but unlike traditional sporadic servers, it does not restore the full budget at every period. Instead, when an aperiodic task arrives, DSS assigns it a deadline and restores only the amount of budget that was actually used. This approach allows the processor to reach full (100 percent) utilization while ensuring that all deadlines are still met. DSS improves the response time of aperiodic tasks without compromising the guarantees of periodic tasks.

Originally introduced by Spuri and Buttazzo[5], DSS has been widely studied for its efficiency in managing mixed task sets under EDF scheduling. This report reviews the theory behind DSS, describes its operation (including budget management), highlights its advantages over traditional methods, and discusses potential application areas.

## I. INTRODUCTION

Real time system often mix periodic tasks (recurring tasks with fixed periods and deadlines) and aperiodic/sporadic tasks (tasks which arrive unpredictably), for example interrupts or user request) And to ensure that the aperiodic tasks receive their service on time without causing problem to any aperiodic tasks to miss their deadlines is a major challenge[2].

A common solution is to use Server Abstraction it is a concept used in real time systems to manage how and when tasks (especially aperiodic or sporadic) are executed. Such as Polling Server, Deferrable server, and sporadic servers. The sporadic servers, introduced by Sprunt et al[4], assigns a fixed priority (typically through Rate Monotonic) to the server and provides a budget  $C_s$  every period  $T_s$ . This allows aperiodic tasks to be executed with limited interference to hard real time tasks. However, under static priority scheduling the processor often cannot be fully utilized, unused server budget might be wasted or caused complex analysis.

With dynamic scheduling (EDF) higher utilization can be achieved. Dynamic priority server as shifted the sporadic server concept to (EDF). In particular, Spuri and Buttazzo proposed the Dynamic Sporadic Server (DSS)[5]. DSS is characterized by a server period  $T_s$  and capacity  $C_s$  (the maximum aperiodic budget)[1]. Unlike SS, DSS does not refill or replenish  $C_s$  fully at each period start. Instead, the server's deadline and replenishment events are set dynamically whenever aperiodic

work arrives or is consumed[5]. This allows DSS to adaptively use idle time for aperiodic tasks and to achieve 100 percent CPU utilization.

But the most important requirement is that under EDF schedule, the sum of the utilization of periodic tasks ( $U_p$ )[1] plus the utilization of DSS ( $U_s = C_s/T_s$ )[1] must satisfy In order to meet all deadlines[5].

$$U_p + U_s \leq 1 \quad [1]$$

Just like EDDF alone, DSS with EDF still allows 100 percent CPU usage, so periodic tasks don't lose processing time. Unlike EDF, a fixed priority sporadic server can't fully use the CPU unless more complex analysis is done[5].

## II. BACKGROUND: PERIODIC AND APERIODIC TASKS; EDF

**Periodic task:** A periodic task  $\tau_i$  generates tasks again and again, each job needs execution time  $C_i$  and must finish by a deadline  $D_i$  after its arrival. commonly,  $D_i = T_i$  (Deadline is equal to Period). The total Utilization of the periodic task is

$$U_p = \sum_i \left( \frac{C_i}{T_i} \right) \quad [1]$$

based on the study by Liu and Layland[3], The EDF algorithm can successfully schedule any set of periodic task (where deadlines are equal to periods) if and only if

$$U_p \leq 1 \quad [1]$$

**Aperiodic tasks:** More general than sporadic, aperiodic tasks include all kind of arrivals, even the ones not allowed under sporadic rules. Aperiodic tasks can be unpredictable, they usually have soft deadlines, which means they can sometimes miss their deadlines. To measure how much CPU they need, we can look at their average or maximum load  $U_a$ . And the main goal of it is to serve these tasks quickly without affecting periodic tasks that have strict deadlines[1]

**sporadic task:** Is a task comes at random times but with a rule there must be minimum gap (period) $T_s$  between two jobs. each job needs some time to run and execute  $C_s$ . Even though the job arrive at irregularly, this minimum gap helps to keep things under control. Because of this sporadic tasks are seen as simpler kind of periodic tasks with more flexible timing.[1].




Fig. 1: Dynamic Sporadic Server example [1]

**EDF Scheduling:** Under EDF, tasks are assigned dynamic priorities equal to their absolute deadline (Earliest deadline = Highest priority). The EDF schedulability criterion rule works with both periodic and aperiodic tasks but only when tasks are treated like sporadic tasks with deadlines[5, 1].

However, simply queuing all periodic tasks under EDF can still cause problems for periodic tasks. Therefore server algorithms are used to control how much CPU the aperiodic workload can use, effectively reserving bandwidth for aperiodic server[1]. DSS is one such server, which is used in EDF based system[1].

### III. SPORADIC AND DYNAMIC SERVERS

#### A. The Sporadic Server (SS)

Before DSS, Sporadic server (SS) was used initially for fixed priority scheduling[1]. The parameters of sporadic server are  $(T_s, C_s)$ . Its rule is When an aperiodic task arrives and the server has available capacity, the server executes that job at high priority as long as it has budget. After consuming a amount of execution  $U$ , the server schedules a refill of  $U$  after one period  $T_s$ [1]. The SS makes sure that the server uses as much of the  $C_s$  time in a period of  $T_s$ , ensuring periodic tasks guarantees only if[1]

$$U_p + \frac{C_s}{T_s} \leq 1 \quad [1]$$

Where  $U_p$  is the utilization of periodic tasks. However, in SS the server priority is fixed, making it not directly suitable for EDF system[1].

#### B. The Dynamic Sporadic Server (DSS)

The Dynamic Sporadic server (DSS) modifies the sporadic server (SS) for EDF Scheduling[5]. Like SS, DSS has a server period  $T_s$  and a Capacity  $C_s$ , DSS has dynamic deadlines

rather than a fixed priority. the server the server behaves like EDF task whose deadline is recalculated at runtime when ever it begins servicing an aperiodic task. its budget does not reset simply at  $T_s$ , instead its only consumed portions are recharged after  $T_s$  units[5, 1].

The DSS behaves in the following way, **Initialization:** The server starts with full budget  $C_s$ . No deadline  $d_s$  is set until the first aperiodic job arrives[5, 1].

**Activation:** When an aperiodic task arrives at time  $t_A$  and the server is idle with available budget  $> 0$ , the server sets its current deadline  $d_s = t_A + T_s$ [5] and schedules its next refill at  $R_T = d_s$ [5]. The server immediately becomes the highest priority EDF task and begins executing aperiodic task. All aperiodic tasks in that busy interval share the same deadline  $d_s$ [5].

**Execution:** The server executes, deducting from its remaining capacity as it runs aperiodic tasks. If the aperiodic workload finishes before the budget is exhausted, the server waits with leftover budget. If more tasks arrive while budget remains, they are queued but served under the same deadline[5].

**Budget Exhaustion or Job Completion:** When capacity is exhausted or the last pending task finishes at time  $t_I$ , the server computes how much budget  $u$  was consumed since  $t_A$  and schedules a refill of amount  $u$  at time  $R_t$ . If pending tasks remain after exhaustion, they are paused until the refill[4, 5].

**Next Activation:** After a refill at every  $R_T$ , if there are tasks pending the server repeats the activation step, it sets a new deadline  $d_s = R_T + T_s$  and executes again.

These rules ensure it never consumes more than the  $C_s$  time in any period  $T_s$ [5, 1]. Thus DSS behaves like aperiodic task under EDF. The schedulability condition for mixing DSS




Fig. 2: Computational demand of a periodic task in  $[t_1, t_2]$  [1].

with periodic task is[1]

$$U_p + U_s = \sum_i \frac{C_i}{T_i} + \frac{C_s}{T_s} \leq 1 \quad [1, 5]$$

Where  $U_p$  is the utilization of periodic task and  $U_s = C_s/T_s$  is servers utilization[1].

**Mathematical Formulation:** Formally, it can be shown (Lemma and Theorem in [5]) that in any busy interval of length  $\Delta$ , the server executes at most  $C_s$  time per  $T_s$  window. Thus, DSS achieves full utilization in EDF-based systems, as long as  $U_p + U_s \leq 1$  [1].

This means the dynamic sporadic server behaves like a periodic task when it comes to CPU usage any interval of the length equal to its period  $T_s$ , it will neever see more than its budget  $C_s$ . so if the utilization from all the periodic tasks plus the DSS remains at or below 100 percent all dadlines can be met. this schedulability condition has been formally proven to be both necessary and sufficient for EDF system using DSS [1, 5].

#### IV. DYNAMIC SPORADIC SERVER SIMULATION

From the Fig 1 Dynamic sporadic example from [1], it shows two periodic tasks  $\tau_1$  and  $\tau_2$  along with a Dynamic Sporadic Server (DSS) that Will work on serving incoming aperiodic requests. task  $\tau_1$  has period  $P_1 = 8$  and execution time  $C_1 = 2$ , while  $\tau_2$  has  $P_2 = 12$  and  $C_2 = 3$ . The Dynamic Sporadic Server (DSS) has a period  $T_s = 6$  and capacity  $C_s = 3$ . aperiodic tasks arrives at spacific times (for example at  $t = 3, 6, 9, \dots$  with different execution demands) and are inlined for service by DSS. The scheduler uses earliest deadline first (EDF) policy. [1].

the DSS logic follows the standard rules: its budget is initialized to  $C_s$  and is refilled one period after being activated for eachtime[1]. whenever a new aperiodic requiest arrives and the server is idle with remaining budget, the server becomes activated, and its deadline is set to (current time +  $T_s$ )[1] while active, the server executes as long as there is pending apeperiodic work and budget, each execution unit decrements both the server's remaining capacity and activity aperiodic job's remaing work. and if the server finises it sbudget and complet all iofits pending tasks, aperiodic tasks. it becomes inactive after that and schedules it s consumed budget to be refilled at the previously assigned deadline [1].

#### V. DSS-BASED UPPAAL MODEL SIMULATION

In this seminar, I'am presenting a dynamic sporadic server (DSS) simulation using UPPAAL. the main goal was to

show how DSS scheduling works and handels aperiodic tasks without intrupting and disturbing the Periodic tasks, which is very important in real times systems. The simulation i made includes two pperiodic tasks, two aperiodic tasks, DSS server, and a controller which acts as the envirnment.

the periodic task are designed tp run at fixed intervals. **PT1** has a period of 8 time units and executes for 2 time unites, while **PT2** has a shorter period 5 unites and it and executes for 1 time unit. Each periodic tasks move through four states: *Idle*, *Ready*, *Execution*, and *Done*. When the period expires, the task moves to **READY**, then executes for its set time, and then goes to execute **DONE** before resetting.

for Aperiodic part, I again included two tasks that are not triggered by time but by events. **APT1** runs for 2 unites and is triggered by the controller at time  $t = 3$ . **APT2** runs for 3 unites and is triggered at  $t = 12$ . however, these tasks can't just start in their own. they must be allowed by the DSS server through a signal called **dss\_grant**. this means that even if a request comes in, the task will only run if the DSS says it can.

The **DSS\_Server** is the key part of the whole model. It manages when and aperiodic task can execute. The capacity of the server is  $C_s$  of 3 units, which means it can be spend up to 3 units on aperiodic task before it refill. the period of replenishment period  $T_s$  is 6 units. the DSS goes through sevrall states, *Idle* (watching for new requests), *Ready* (when it decides to grant access), *Execution* (running an aperiodic task), *Done* (after the task finishes), and *WaitReplenish* (when it has no capacity and is waiting to refill). It keeps track of how much capacity it has and also remembers when to refill again. with **d\_server**.

The **Controller** in the model is a simple component that simulates the enviourment. It keeps track of time and triggers the aperiodic task requests at specific points. in my simulation, it sets **aperiodic\_request1 = true** at  $t = 3$  and **aperiodic\_request2 = true** at  $t = 12$ .

When the simulator runs, both periodic tasks start in **Idle**. at time  $t = 3$ , When the first aperiodic task arives it is then triggered by the controller. scince DSS has full capacity at the start, it grants permission, and **AperiodicTask1** execute, then the first replenishment time and the server is. Scince  $d_s$  is the earliest deadline first, the highest priority shifts toward DSS task in th esystem and the request gets served unti completion[1]. This reduces the the servers capacity from 3 time units to 2 units then again the periodic tasks continue executing according to their periods **PeriodicTask2** runs at  $t = 5$  and **PeriodicTask1** at  $t = 8$ . At  $t = 12$ , the controller sends the second aperiodic request[1]. Now depending how much capacity the DSS has left, it will run the second task for that period and after he refil it again runs it[1].

The simulation clearly shows how DSS works and how it serves the aperiodic tasks without intrupting the periodic tasks. The DSS carefully tracks how much capacity is used and refills only after the defined period. This method ensures that periodic tasks always get priority and deadlines are not missed, while still giving space to aperiodic tasks when possible. I believe this model demonstrates the concept in a simple but

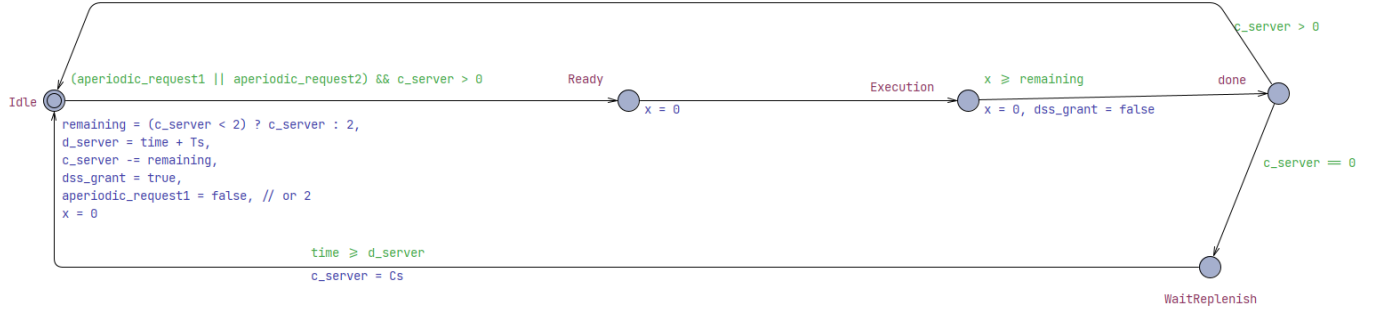


Fig. 3: DSS server




Fig. 4: Controller

effective way, and helps to understand real time scheduling more practically.

## VI. COMPARISON WITH OTHER SERVERS

Table I shows the comparsion of the key features of Polling,Deferrable,sporadic and dynamic sporadic server. polling server operate under fixed priority scheduling and re-sets it budget after each period regardless of use [4]. Deferrable servers also run at high prority but carry their unused capacit till the end of the periode [1]. The classic Sporadic server fro Ratemonotonic schedullng only replenishes only the portion of their buddget that was actually used [4].

Because of the EDF and its flexible budget usage Dynamic sporadic server stands apart from others. In practice, DSS provides the responsiveness of a high priority server without leaving CPU time idle when aperiodic demand is present [1, 2].

## VII. APPLICATIONS AND USE CASES

Dynamic Sporadic Servers (DSS) are well-suited to systems that must mix hard real-time periodic tasks with unpredictable aperiodic workloads. Typical examples include automotive ECUs, robotic controllers, avionics systems, embedded Linux platforms, multimedia devices, and medical equipment.

For instance, an automotive braking controller may have a periodic speed-monitoring task and a sporadic braking-request task [5]. In such domains, DSS provides a reserved capacity

Feature	Polling	Deferrable	Sporadic	DSS
Scheduling	Fixed-priority (RM)	Fixed-priority (RM)	Fixed-priority (RM)	Dynamic (EDF)
Priority	Highest static	Highest static	Static (RM)	Dynamic (deadline $t + T_s$ )
Replenishment	Periodic full	Periodic full	On-demand	On-demand
Unused Budget	Lost if unused	Carried to period end	Carried to next period	Reused flexibly
Utilization	Often sub-optimal	Improved response	Moderate	Full (up to 100%)

TABLE I: Comparison of Aperiodic Servers

for aperiodic jobs without violating hard deadlines of periodic jobs.

In robotics, a servo loop may run at a fixed rate while sensor or vision processing tasks arrive unpredictably. DSS can absorb these aperiodic jobs using its time-varying deadline mechanism (e.g., “ties are always resolved in favor of the server”) [1, 5], improving responsiveness without jeopardizing safety-critical loops.

Similarly, avionics flight computers often interleave periodic control tasks (e.g., stabilization) with aperiodic ones

(e.g., communication, diagnostics). In multimedia systems, the Linux `SCHED_DEADLINE` scheduler uses a variant of DSS (Constant Bandwidth Server) to serve aperiodic workloads (network packets, user interactions) alongside real-time tasks [1, 5].

In medical systems such as implantable monitors or patient alarms, DSS can isolate life-critical periodic sensing from sporadic events like logging or alarms.

- **Automotive ECUs:** DSS allocates a fixed budget  $C_s$  per period  $T_s$  to event-driven aperiodic tasks without delaying periodic engine control jobs [5].
- **Robotics and Control:** DSS enables mixing real-time servo loops with irregular processing (vision/path planning) while meeting EDF feasibility conditions [1].
- **Avionics:** DSS helps meet deadlines for both periodic sensor/actuator cycles and asynchronous diagnostics.
- **Embedded Linux & Multimedia:** DSS variants like CBS are used in Linux for EDF-based scheduling of sporadic user interactions or audio/video streams [1].
- **Medical Systems:** DSS ensures timely response to alarms while maintaining guaranteed periodic monitoring rates.

DSS offers strong temporal isolation: aperiodic tasks are treated like a periodic task with utilization  $U_s = C_s/T_s$ . Spuri and Buttazzo showed that if  $U_p + U_s \leq 1$ , all EDF deadlines will be met [5, 1].

#### VIII. LIMITATIONS AND PRACTICAL CONSIDERATIONS

Implementing DSS involves non-trivial complexity. The operating system must manage a dynamic deadline and a replenishable budget. Each aperiodic job arrival or execution requires updating the deadline and tracking consumed capacity [1, 5]. This includes scheduling a replenishment for any consumed portion after  $T_s$  time units.

Whenever an aperiodic job executes, it inherits the server's deadline and may preempt periodic jobs. This makes DSS more complex than simple polling or deferrable servers [5, 2, 1].

In systems using fixed-priority scheduling (e.g., Rate Monotonic), DSS cannot be directly implemented. It relies on EDF for dynamic prioritization. While real-time Linux and some RTOSes support CBS/DSS, integration still requires modifying scheduler internals.

There is also an overhead trade-off: smaller  $T_s$  reduces latency but increases context switching. A short  $T_s$  means smaller budgets  $C_s$ , which may lead to queue buildup. Conversely, large  $T_s$  values reduce overhead but delay service to aperiodic jobs [5, 2].

Heavy aperiodic workloads can overwhelm the server. If job arrivals exceed  $U_s$ , delays increase sharply. Worst-case latency occurs when jobs arrive just after budget exhaustion. System designers must choose  $T_s$  and  $C_s$  carefully to balance responsiveness with feasibility.

Corner cases also matter. If a server and a periodic job become ready at the same time, tie-breaking policies must prioritize the server [1, 5]. Priority inheritance protocols are needed to avoid deadlocks when tasks share resources. In

multicore systems, sharing the server budget across cores further complicates scheduling.

In summary, DSS offers predictability but requires careful implementation and tuning.

#### IX. CONCLUSION

The Dynamic Sporadic Server is a powerful scheduling mechanism that extends EDF by offering guaranteed bandwidth to aperiodic tasks. By dynamically assigning deadlines and replenishing only the used capacity, DSS achieves efficient processor utilization while respecting hard deadlines for periodic tasks [5, 1].

Its strengths include:

- Strong temporal isolation and bandwidth reservation.
- High responsiveness to aperiodic jobs.
- Full utilization up to 100% under EDF schedulability bounds.

DSS is particularly beneficial in mixed-criticality real-time systems such as automotive, avionics, robotics, and embedded Linux environments. Compared to simpler servers, it provides better control and responsiveness, especially under unpredictable workloads.

Looking forward, DSS can be enhanced through variants like the Constant Bandwidth Server (CBS), integration with priority-inheritance protocols, and adaptations for multicore platforms [2]. Future research may also explore hybrid EDF-fixed priority schemes and virtualization-aware servers.

In conclusion, DSS remains a relevant and effective solution for real-time systems that require dynamic, predictable aperiodic task management

#### REFERENCES

- [1] Giorgio C. Buttazzo. *Hard Real-Time Computing Systems: Predictable Scheduling Algorithms and Applications*. Springer Science & Business Media, 2011.
- [2] Philip A. Laplante. *Real-Time Systems Design and Analysis: Tools for the Practitioner*. 4th. Wiley-IEEE Press, 2011.
- [3] C. L. Liu and J. W. Layland. "Scheduling algorithms for multiprogramming in a hard-real-time environment". In: *Journal of the ACM (JACM)* 20.1 (1973), pp. 46–61.
- [4] Barry Sprunt, Lui Sha, and John P. Lehoczky. "Aperiodic Task Scheduling for Hard-Real-Time Systems". In: *The Journal of Real-Time Systems* 1.1 (1989), pp. 27–60.
- [5] Marco Spuri and Giorgio C. Buttazzo. "Efficient Aperiodic Service under Earliest Deadline Scheduling". In: *Proceedings of the IEEE Real-Time Systems Symposium (RTSS)*. IEEE, 1994, pp. 2–11.