

Dynamic Sporadic Server

Moiz Zaheer Malik

B.Eng. Electronic Engineering

Hochschule Hamm-Lippstadt

Lippstadt, Germany

moiz-zaheer.malik@stud.hshl.de

Abstract—In real-time systems, where tasks have different levels of critical importance, it is essential to serve aperiodic (irregular, event-driven) tasks while ensuring that the deadlines of high-priority periodic tasks are not violated. The Dynamic Sporadic Server (DSS) is a scheduling method designed for Earliest Deadline First (EDF) systems that addresses this problem.

DSS is defined by a period T_s and a budget C_s , but unlike traditional sporadic servers, it does not restore the full budget at every period. Instead, when an aperiodic task arrives, DSS assigns it a deadline and restores only the amount of budget that was actually used. This approach allows the processor to reach full (100 percent) utilization while ensuring that all deadlines are still met. DSS improves the response time of aperiodic tasks without compromising the guarantees of periodic tasks.

Originally introduced by Spuri and Buttazzo[5], DSS has been widely studied for its efficiency in managing mixed task sets under EDF scheduling. This report reviews the theory behind DSS, describes its operation (including budget management), highlights its advantages over traditional methods, and discusses potential application areas.

I. INTRODUCTION

Real time system often mix periodic tasks (recurring tasks with fixed periods and deadlines) and aperiodic/sporadic tasks (tasks which arrive unpredictably), for example interrupts or user request) And to ensure that the aperiodic tasks receive their service on time without causing problem to any aperiodic tasks to miss their deadlines is a major challenge[2].

A common solution is to use Server Abstraction it is a concept used in real time systems to manage how and when tasks (especially aperiodic or sporadic) are executed. Such as Polling Server, Deferrable server, and sporadic servers. The sporadic servers, introduced by Sprunt et al[4], assigns a fixed priority (typically through Rate Monotonic) to the server and provides a budget C_s every period T_s . This allows aperiodic tasks to be executed with limited interference to hard real time tasks. However, under static priority scheduling the processor often cannot be fully utilized, unused server budget might be wasted or caused complex analysis.

With dynamic scheduling (EDF) higher utilization can be achieved. Dynamic priority server as shifted the sporadic server concept to (EDF). In particular, Spuri and Buttazzo proposed the Dynamic Sporadic Server (DSS)[5]. DSS is characterized by a server period T_s and capacity C_s (the maximum aperiodic budget)[1]. Unlike SS, DSS does not refill or replenish C_s fully at each period start. Instead, the server's deadline and replenishment events are set dynamically whenever aperiodic

work arrives or is consumed[5]. This allows DSS to adaptively use idle time for aperiodic tasks and to achieve 100 percent CPU utilization.

But the most important requirement is that under EDF schedule, the sum of the utilization of periodic tasks (U_p)[1] plus the utilization of DSS ($U_s = C_s/T_s$)[1] must satisfy In order to meet all deadlines[5].

$$U_p + U_s \leq 1 \quad [1]$$

Just like EDDF alone, DSS with EDF still allows 100 percent CPU usage, so periodic tasks don't lose processing time. Unlike EDF, a fixed priority sporadic server can't fully use the CPU unless more complex analysis is done[5].

II. BACKGROUND: PERIODIC AND APERIODIC TASKS; EDF

Periodic task: A periodic task τ_i generates tasks again and again, each job needs execution time C_i and must finish by a deadline D_i after its arrival. commonly, $D_i = T_i$ (Deadline is equal to Period). The total Utilization of the periodic task is

$$U_p = \sum_i \left(\frac{C_i}{T_i} \right) \quad [1]$$

based on the study by Liu and Layland[3], The EDF algorithm can successfully schedule any set of periodic task (where deadlines are equal to periods) if and only if

$$U_p \leq 1 \quad [1]$$

Aperiodic tasks: More general than sporadic, aperiodic tasks include all kind of arrivals, even the ones not allowed under sporadic rules. Aperiodic tasks can be unpredictable, they usually have soft deadlines, which means they can sometimes miss their deadlines. To measure how much CPU they need, we can look at their average or maximum load U_a . And the main goal of it is to serve these tasks quickly without affecting periodic tasks that have strict deadlines[1]

sporadic task: Is a task comes at random times but with a rule there must be minimum gap (period) T_s between two jobs. each job needs some time to run and execute C_s . Even though the job arrive at irregularly, this minimum gap helps to keep things under control. Because of this sporadic tasks are seen as simpler kind of periodic tasks with more flexible timing.[1].

EDF Scheduling: Under EDF, tasks are assigned dynamic priorities equal to their absolute deadline (Earliest deadline = Highest priority). The EDF schedulability criterion rule works with both periodic and aperiodic tasks but only when tasks are treated like sporadic tasks with deadlines[5, 1].

However, simply queuing all periodic tasks under EDF can still cause problems for periodic tasks. Therefore server algorithms are used to control how much CPU the aperiodic workload can use, effectively reserving bandwidth for aperiodic server[1]. DSS is one such server, which is used in EDF based system[1].

III. SPORADIC AND DYNAMIC SERVERS

A. The Sporadic Server (SS)

Before DSS, Sporadic server (SS) was used initially for fixed priority scheduling[1]. The parameters of sporadic server are (T_s, C_s) . Its rule is When an aperiodic task arrives and the server has available capacity, the server executes that job at high priority as long as it has budget. After consuming a amount of execution U , the server schedules a refill of U after one period T_s [1]. The SS makes sure that the server uses as much of the C_s time in a period of T_s , ensuring periodic tasks guarantees only if[1]

$$U_p + \frac{C_s}{T_s} \leq 1 \quad [1]$$

Where U_p is the utilization of periodic tasks. However, in SS the server priority is fixed, making it not directly suitable for EDF system[1].

B. The Dynamic Sporadic Server (DSS)

The Dynamic Sporadic server (DSS) modifies the sporadic server (SS) for EDF Scheduling[5]. Like SS, DSS has a server period T_s and a Capacity C_s , DSS has dynamic deadlines rather than a fixed priority. The server the server behaves like EDF task whose deadline is recalculated at runtime whenever it begins servicing an aperiodic task. Its budget does not reset simply at T_s , instead its only consumed portions are recharged after T_s units[5, 1].

The DSS behaves in the following way, **Initialization:** The server starts with full budget C_s . No deadline d_s is set until the first aperiodic job arrives[5, 1].

Activation: When an aperiodic task arrives at time t_A and the server is idle with available budget > 0 , the server sets its current deadline $d_s = t_A + T_s$ [5] and schedules its next refill at $R_T = d_s$ [5]. The server immediately becomes the highest priority EDF task and begins executing aperiodic task. All aperiodic tasks in that busy interval share the same deadline d_s [5].

Execution: The server executes, deducting from its remaining capacity as it runs aperiodic tasks. If the aperiodic workload finishes before the budget is exhausted, the server waits with leftover budget. If more tasks arrive while budget remains, they are queued but served under the same deadline[5].

Budget Exhaustion or Job Completion: When capacity is exhausted or the last pending task finishes at time t_I , the




Fig. 1: Dynamic Sporadic Server example[1]

server computes how much budget u was consumed since t_A and schedules a refill of amount u at time R_t . If pending tasks remain after exhaustion, they are paused until the refill[4, 5].

Next Activation: After a refill at every R_T , if there are tasks pending the server repeats the activation step, it sets a new deadline $d_s = R_T + T_s$ and executes again.

These rules ensure it never consumes more than the C_s time in any period T_s [5, 1]. Thus DSS behaves like an aperiodic task under EDF. The schedulability condition for mixing DSS with periodic task is[1]

$$U_p + U_s = \sum_i \frac{C_i}{T_i} + \frac{C_s}{T_s} \leq 1 \quad [1, 5]$$

Where U_p is the utilization of periodic task and $U_s = C_s/T_s$ is server utilization[1].

Mathematical Formulation: Formally, it can be shown (Lemma and Theorem in [5]) that in any busy interval of length Δ , the server executes at most C_s time per T_s window. Thus, DSS achieves full utilization in EDF-based systems, as long as $U_p + U_s \leq 1$ [1].

IV. DYNAMIC SPORADIC SERVER SIMULATION

In this example, we simulate two periodic tasks 1 and 2 along with a Dynamic Sporadic Server (DSS) that serves incoming aperiodic requests. Task 1 has period $P_1 = 8$ and execution time $C_1 = 2$, while 2 has $P_2 = 12$ and $C_2 = 3$. The DSS is configured with period $T_s = 6$ and capacity $C_s = 3$. Aperiodic tasks arrive at specific times (for example at $t = 3, 6, 9, \dots$ with varying execution demands) and are queued for service by the DSS. The scheduler uses earliest-deadline-first (EDF) policy, breaking ties in favor of the server[2, 8]. The DSS logic follows the standard rules: its budget is initialized to C_s and is replenished one period after each activation[8]. Whenever a new aperiodic request arrives and the server is idle with remaining budget, the server becomes active, and its deadline is set to $(\text{current time} + T_s)$ [8]. While active, the server executes as long as there is pending aperiodic work and budget remains; each execution unit decrements both the server's remaining capacity and the active aperiodic job's remaining work. If the server exhausts its budget or completes all queued aperiodic tasks, it becomes inactive




Fig. 2: Gantt chart

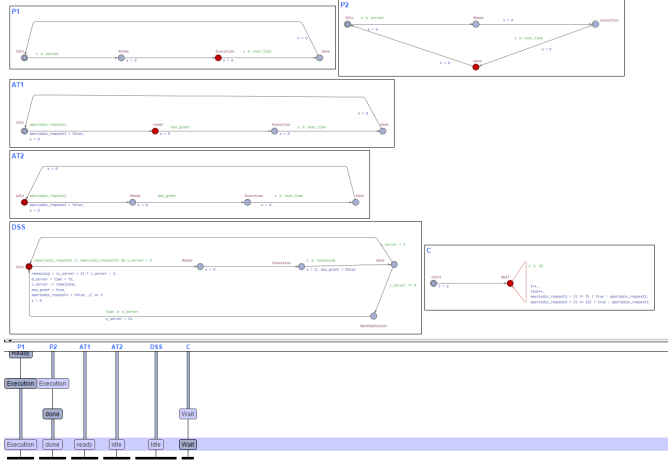


Fig. 3: Symbolic Simulator

and schedules its consumed budget to be replenished at the previously assigned deadline[1]

V. DSS-BASED UPPAAL MODEL SIMULATION

In this seminar, I presented a simulation of the Dynamic Sporadic Server (DSS) model using the UPPAAL tool. The main aim was to show how DSS scheduling can handle aperiodic tasks without interfering with periodic tasks, which is very important in real-time systems. The simulation I made




Fig. 4: DSS server




Fig. 5: Controller



Fig. 6: Periodic Task 1




Fig. 7: Periodic Task 2

includes two periodic tasks, two aperiodic tasks, a DSS server, and a controller that acts as the environment.

The periodic tasks are designed to run at fixed intervals. **PeriodicTask1** has a period of 8 time units and executes for 2 time units, while **PeriodicTask2** has a shorter period of 5 units and executes for 1 time unit. Each periodic task moves through four states: *Idle*, *Ready*, *Execution*, and *Done*. When the period expires, the task moves to *Ready*, then executes for its set time, and then goes to *Done* before resetting.

For the aperiodic part, I included two tasks that are not triggered by time but by events. **AperiodicTask1** runs for 2 units and is triggered by the controller at time $t = 3$. **AperiodicTask2** runs for 3 units and is triggered at $t = 12$. However, these tasks can't just start on their own. They must be allowed by the DSS server through a signal called *dss_grant*. This means that even if a request comes in, the task will only run if the DSS says it can.

The **DSS_Server** is the key part of the whole model. It manages when and how aperiodic tasks can execute. The server has a capacity (C_s) of 3 units, which means it can spend up to 3 time units on aperiodic tasks before it needs to refill. The replenishment period (T_s) is 6 units, so every 6 units, it gets its full capacity back. The DSS goes through several states: *Idle* (watching for new requests), *Ready* (when it decides to grant access), *Execution* (running an aperiodic task), *Done* (after the task finishes), and *WaitReplenish* (when it has no capacity and is waiting to refill). It keeps track of how much capacity it has left using the variable *c_server* and also remembers when to refill with *d_server*.

The **Controller** in the model is a simple component that



Fig. 8: Aperiodic Task 1



Fig. 9: Aperiodic Task 2

simulates the environment. It keeps track of time and triggers the aperiodic task requests at specific points. In my simulation, it sets `aperiodic_request1 = true` at $t = 3$ and `aperiodic_request2 = true` at $t = 12$. It also increases logical time step by step.

When the simulation runs, both periodic tasks start in *Idle*. At time $t = 3$, the controller triggers the first aperiodic request. Since DSS has full capacity at the start, it grants permission, and **AperiodicTask1** executes. This reduces the server’s capacity by 2 units. Then, periodic tasks keep executing according to their periods **PeriodicTask2** runs at $t = 5$ and **PeriodicTask1** at $t = 8$. At $t = 12$, the controller sends the second aperiodic request. Now, depending on how much capacity the DSS has left, it may either run the task or wait until its capacity is replenished.

The simulation clearly shows how DSS works to let aperiodic tasks execute without breaking the schedule of the periodic ones. The DSS carefully tracks how much capacity is used and refills only after the defined period. This method ensures that periodic tasks always get priority and deadlines are not missed, while still giving space to aperiodic tasks when possible. I believe this model demonstrates the concept in a simple but effective way, and helps to understand real-time scheduling more practically.

VI. COMPARISON WITH OTHER SERVERS

Table I compares key features of Polling, Deferrable, Sporadic, and Dynamic Sporadic servers. Polling servers operate under fixed-priority scheduling and reset their budget at each period regardless of use [4]. Deferrable servers also run at high priority but carry unused capacity to the end of the period [1]. The classic Sporadic Server (for Rate-Monotonic scheduling) replenishes only the portion of its budget that was actually consumed [4].

Feature	Polling	Deferrable	Sporadic	DSS
Scheduling	Fixed-priority (RM)	Fixed-priority (RM)	Fixed-priority (RM)	Dynamic (EDF)
Priority	Highest static	Highest static	Static (RM)	Dynamic (deadline $t + T_s$)
Replenishment	Periodic full	Periodic full	On-demand	On-demand
Unused Budget	Lost if unused	Carried to period end	Carried to next period	Reused flexibly
Utilization	Often sub-optimal	Improved response	Moderate	Full (up to 100%)

TABLE I: Comparison of Aperiodic Servers

DSS stands apart by its EDF-based policy and flexible budget usage. In practice, DSS provides the responsiveness of a high-priority server without leaving CPU time idle when aperiodic demand is present [1, 2].

VII. APPLICATIONS AND USE CASES

REFERENCES

- [1] Giorgio C. Buttazzo. *Hard Real-Time Computing Systems: Predictable Scheduling Algorithms and Applications*. Springer Science & Business Media, 2011.
- [2] Philip A. Laplante. *Real-Time Systems Design and Analysis: Tools for the Practitioner*. 4th. Wiley-IEEE Press, 2011.
- [3] C. L. Liu and J. W. Layland. “Scheduling algorithms for multiprogramming in a hard-real-time environment”. In: *Journal of the ACM (JACM)* 20.1 (1973), pp. 46–61.
- [4] Barry Sprunt, Lui Sha, and John P. Lehoczky. “Aperiodic Task Scheduling for Hard-Real-Time Systems”. In: *The Journal of Real-Time Systems* 1.1 (1989), pp. 27–60.
- [5] Marco Spuri and Giorgio C. Buttazzo. “Efficient Aperiodic Service under Earliest Deadline Scheduling”. In: *Proceedings of the IEEE Real-Time Systems Symposium (RTSS)*. IEEE, 1994, pp. 2–11.